

ilearn-healthcare-capstone-project

August 14, 2023

0.0.1 Importing Libraries

```
[203]: %matplotlib inline
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
```

0.0.2 Loading Dataset

```
[204]: data = pd.read_csv('health_care_diabetes.csv')
```

```
[205]: data.head()
```

```
[205]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[206]: data.shape
```

```
[206]: (768, 9)
```

0.1 Project Task: Week 1 – Data Exploration and Missing Values Treatment

```
[304]: #Checking for null values in Dataset
data.isnull().any()
```

```
[304]: Pregnancies      False
        Glucose         False
        BloodPressure   False
        SkinThickness   False
        Insulin         False
        BMI             False
        DiabetesPedigreeFunction False
        Age            False
        Outcome         False
        dtype: bool
```

Since the 0 value in Glucose,BloodPressure,SkinThickness,Insulin and BMI variables represent missing values.Lets find now many instances are there in each of the above variables

```
[208]: data[data['Glucose']==0]
```

```
[208]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
75	1	0	48	20	0	24.7	
182	1	0	74	20	23	27.7	
342	1	0	68	35	0	32.0	
349	5	0	80	32	0	41.0	
502	6	0	68	41	0	39.0	

	DiabetesPedigreeFunction	Age	Outcome
75	0.140	22	0
182	0.299	21	0
342	0.389	22	0
349	0.346	37	1
502	0.727	41	1

```
[209]: (5/765)*100
        #only 0.6% of data is having missing values in Glucose column. No need to worry
        ↳we can ignore them
```

```
[209]: 0.6535947712418301
```

```
[210]: (data[data['BloodPressure']==0]).shape
```

```
[210]: (35, 9)
```

```
[211]: (35/765)*100
        #4.5% of data is having missing values in BloodPressure column
```

```
[211]: 4.57516339869281
```

```
[212]: (data[data['SkinThickness']==0]).shape
```

```
[212]: (227, 9)
```

```
[213]: (227/765)*100  
#29.6% of data is having missing values in SkinThickness column
```

```
[213]: 29.673202614379086
```

```
[214]: (data[data['Insulin']==0]).shape
```

```
[214]: (374, 9)
```

```
[215]: (374/765)*100  
#~49% of data is having missing values in Insulin column
```

```
[215]: 48.888888888888886
```

```
[216]: (data[data['BMI']==0]).shape
```

```
[216]: (11, 9)
```


```
[217]: (11/765)*100  
#1.4% of data is having missing values in BMI column
```

```
[217]: 1.4379084967320261
```

Since Insulin and SkinThickness are having higher percentages of missing values lets try to fill up the missing values

```
[218]: plt.hist(data['SkinThickness'],edgecolor='red')
```

```
[218]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),  
      array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),  
      <BarContainer object of 10 artists>)
```




```
[219]: data[data['SkinThickness']!=0]['SkinThickness'].describe()
```

```
[219]: count      541.000000
      mean       29.153420
      std        10.476982
      min         7.000000
      25%        22.000000
      50%        29.000000
      75%        36.000000
      max        99.000000
      Name: SkinThickness, dtype: float64
```

```
[220]: plt.hist(data['Insulin'],edgecolor='red')
```

```
[220]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,
              761.4, 846. ]),
      <BarContainer object of 10 artists>)
```



```
[221]: data[data['Insulin']!=0]['Insulin'].describe()
```

```
[221]: count    394.000000
      mean     155.548223
      std      118.775855
      min       14.000000
      25%       76.250000
      50%      125.000000
      75%      190.000000
      max      846.000000
      Name: Insulin, dtype: float64
```


Mean value of Skinthickness is ~29 and the mean value of Insulin is ~155 let impute the missing values with means

```
[222]: from numpy import nan
      dataset_imputed = data
      dataset_imputed[['SkinThickness','Insulin']] =_
      ↪dataset_imputed[['SkinThickness','Insulin']].replace(0, nan)
```

```
[223]: dataset_imputed.fillna(dataset_imputed.mean(), inplace=True)
```

```
[224]: plt.hist(dataset_imputed['Insulin'],edgecolor='red')
```

```
[224]: (array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,
       762.8, 846. ]),
<BarContainer object of 10 artists>)
```



```
[225]: data.describe()
```

```
[225]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	29.153420	155.548223
std	3.369578	31.972618	19.355807	8.790942	85.021108
min	0.000000	0.000000	0.000000	7.000000	14.000000
25%	1.000000	99.000000	62.000000	25.000000	121.500000
50%	3.000000	117.000000	72.000000	29.153420	155.548223
75%	6.000000	140.250000	80.000000	32.000000	155.548223
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000

max	67.100000	2.420000	81.000000	1.000000
-----	-----------	----------	-----------	----------

```
[226]: dataset_imputed.describe()
```

```
[226]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	29.153420	155.548223
std	3.369578	31.972618	19.355807	8.790942	85.021108
min	0.000000	0.000000	0.000000	7.000000	14.000000
25%	1.000000	99.000000	62.000000	25.000000	121.500000
50%	3.000000	117.000000	72.000000	29.153420	155.548223
75%	6.000000	140.250000	80.000000	32.000000	155.548223
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[227]: dataset_imputed.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    float64
4   Insulin                             768 non-null    float64
5   BMI                                 768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(4), int64(5)
memory usage: 54.1 KB
```

```
[228]: Positive = dataset_imputed[dataset_imputed['Outcome']==1]
Positive.head(5)
```

```
[228]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35.00000	155.548223	33.6	
2	8	183	64	29.15342	155.548223	23.3	
4	0	137	40	35.00000	168.000000	43.1	
6	3	78	50	32.00000	88.000000	31.0	
8	2	197	70	45.00000	543.000000	30.5	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
2	0.672	32	1
4	2.288	33	1
6	0.248	26	1
8	0.158	53	1

```
[229]: Negative = dataset_imputed[dataset_imputed['Outcome']==0]
Negative.head(5)
```

```
[229]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
1	1	85	66	29.00000	155.548223	26.6	
3	1	89	66	23.00000	94.000000	28.1	
5	5	116	74	29.15342	155.548223	25.6	
7	10	115	0	29.15342	155.548223	35.3	
10	4	110	92	29.15342	155.548223	37.6	


	DiabetesPedigreeFunction	Age	Outcome
1	0.351	31	0
3	0.167	21	0
5	0.201	30	0
7	0.134	29	0
10	0.191	30	0

```
[230]: dataset_imputed['Glucose'].value_counts().head(5)
```

```
[230]: 99      17
100     17
111     14
129     14
125     14
Name: Glucose, dtype: int64
```

```
[231]: plt.hist(dataset_imputed['Glucose'],edgecolor='red')
```

```
[231]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
array([ 0. , 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
179.1, 199. ]),
<BarContainer object of 10 artists>)
```





```
[232]: dataset_imputed['BloodPressure'].value_counts().head(7)
```

```
[232]: 70      57
       74      52
       78      45
       68      45
       72      44
       64      43
       80      40
       Name: BloodPressure, dtype: int64
```

```
[233]: plt.hist(dataset_imputed['BloodPressure'],edgecolor='red')
```

```
[233]: (array([ 35.,   1.,   2.,  13., 107., 261., 243.,  87.,  14.,   5.]),
       array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,
              109.8, 122. ]),
       <BarContainer object of 10 artists>)
```




```
[234]: dataset_imputed['SkinThickness'].value_counts().head(7)
```

```
[234]: 29.15342    227
      32.00000     31
      30.00000     27
      27.00000     23
      23.00000     22
      33.00000     20
      28.00000     20
      Name: SkinThickness, dtype: int64
```

```
[235]: plt.hist(dataset_imputed['SkinThickness'],edgecolor='red')
```

```
[235]: (array([ 59., 141., 408., 118., 36., 4., 1., 0., 0., 1.]),
      array([ 7. , 16.2, 25.4, 34.6, 43.8, 53. , 62.2, 71.4, 80.6, 89.8, 99. ]),
      <BarContainer object of 10 artists>)
```




```
[236]: dataset_imputed['Insulin'].value_counts().head(7)
```

```
[236]: 155.548223    374
      105.000000     11
      130.000000      9
      140.000000      9
      120.000000      8
      94.000000       7
      180.000000       7
      Name: Insulin, dtype: int64
```

```
[237]: plt.hist(dataset_imputed['Insulin'],edgecolor='red')
```

```
[237]: (array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
      array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,
              762.8, 846. ]),
      <BarContainer object of 10 artists>)
```




```
[238]: dataset_imputed['BMI'].value_counts().head(7)
```

```
[238]: 32.0    13
      31.6    12
      31.2    12
      0.0    11
      32.4    10
      33.3    10
      30.1     9
      Name: BMI, dtype: int64
```

```
[239]: plt.hist(dataset_imputed['BMI'],edgecolor='red')
```

```
[239]: (array([ 11.,   0.,  15., 156., 268., 224.,  78.,  12.,   3.,   1.]),
      array([ 0.   ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
            60.39, 67.1 ]),
      <BarContainer object of 10 artists>)
```



```
[240]: dataset_imputed.describe().transpose()
```

```
[240]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	29.153420	8.790942	7.000	25.00000	
Insulin	768.0	155.548223	85.021108	14.000	121.50000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.000000	6.000000	17.00
Glucose	117.000000	140.250000	199.00
BloodPressure	72.000000	80.000000	122.00
SkinThickness	29.153420	32.000000	99.00
Insulin	155.548223	155.548223	846.00
BMI	32.000000	36.600000	67.10
DiabetesPedigreeFunction	0.372500	0.626250	2.42
Age	29.000000	41.000000	81.00
Outcome	0.000000	1.000000	1.00

0.1.1 Project Task: Week 2 – Corelation Analysis and Scatter Plots

```
[241]: Positive.shape
```


```
[241]: (268, 9)
```

```
[242]: Negative.shape
```

```
[242]: (500, 9)
```

```
[243]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20,edgecolor='red')
```

```
[243]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
        14.,  7.,  3.,  1.,  1.,  0.,  1.]),
       array([ 0.    ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
        26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
        53.68 , 57.035, 60.39 , 63.745, 67.1  ]),
       [<matplotlib.patches.Polygon at 0x7f93bbaaa640>])
```




```
[244]: Positive['BMI'].value_counts().head(7)
```

```
[244]: 32.9      8
      31.6      7
      33.3      6
      31.2      5
      30.5      5
```

```
32.0    5
34.3    4
Name: BMI, dtype: int64
```

```
[245]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20,edgecolor='red')
```

```
[245]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
       array([ 0.   ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199.   ]),
       [<matplotlib.patches.Polygon at 0x7f93bb3d4490>])
```




```
[246]: Positive['Glucose'].value_counts().head(7)
```

```
[246]: 125    7
      128    6
      129    6
      115    6
      158    6
      146    5
      124    5
      Name: Glucose, dtype: int64
```

```
[247]: plt.
        ↪hist(Positive['BloodPressure'],histtype='stepfilled',bins=20,edgecolor='red')
```

```
[247]: (array([16., 0., 0., 0., 0., 1., 0., 1., 6., 6., 19., 37., 56.,
        36., 41., 31., 7., 4., 4., 3.]),
        array([ 0.,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
        51.3, 57., 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
        102.6, 108.3, 114. ]),
        [<matplotlib.patches.Polygon at 0x7f93d4849e50>])
```




```
[248]: Positive['BloodPressure'].value_counts().head(7)
```

```
[248]: 70    23
        76    18
        78    17
        74    17
        72    16
         0    16
        80    13
        Name: BloodPressure, dtype: int64
```

```
[249]: plt.
        ↪hist(Positive['SkinThickness'],histtype='stepfilled',bins=20,edgecolor='red')
```



```
[249]: (array([ 1.,  5., 11., 21., 113., 41., 34., 20., 15., 4., 1.,
              0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  1.]),
        array([ 7. , 11.6, 16.2, 20.8, 25.4, 30. , 34.6, 39.2, 43.8, 48.4, 53. ,
              57.6, 62.2, 66.8, 71.4, 76. , 80.6, 85.2, 89.8, 94.4, 99. ]),
        [<matplotlib.patches.Polygon at 0x7f93d2eb1940>])
```




```
[250]: Positive['SkinThickness'].value_counts().head(7)
```

```
[250]: 29.15342    88
       32.00000    14
       30.00000     9
       33.00000     9
       39.00000     8
       37.00000     8
       36.00000     8
       Name: SkinThickness, dtype: int64
```

```
[251]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20,edgecolor='red')
```

```
[251]: (array([ 4., 12., 27., 169., 18., 10.,  8.,  5.,  2., 1., 1.,
              6.,  2.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
        array([ 14. , 55.6, 97.2, 138.8, 180.4, 222. , 263.6, 305.2, 346.8,
              388.4, 430. , 471.6, 513.2, 554.8, 596.4, 638. , 679.6, 721.2,
              762.8, 804.4, 846. ]),
        [<matplotlib.patches.Polygon at 0x7f93d449bf40>])
```



```
[252]: Positive['Insulin'].value_counts().head(7)
```

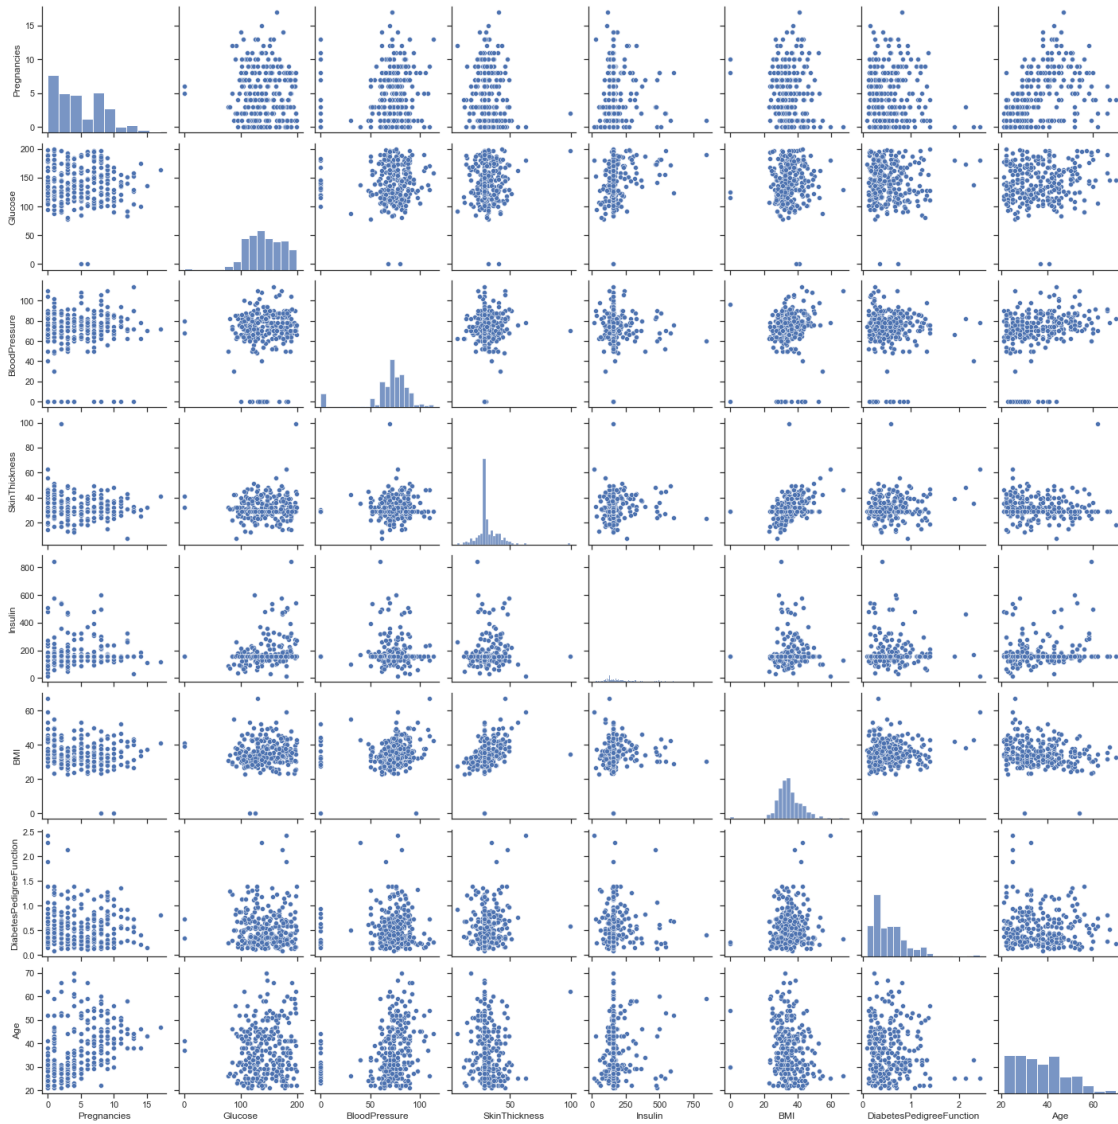
```
[252]: 155.548223    138
      130.000000     6
      180.000000     4
      175.000000     3
      156.000000     3
      185.000000     2
      194.000000     2
      Name: Insulin, dtype: int64
```

0.1.2 Scatter plots

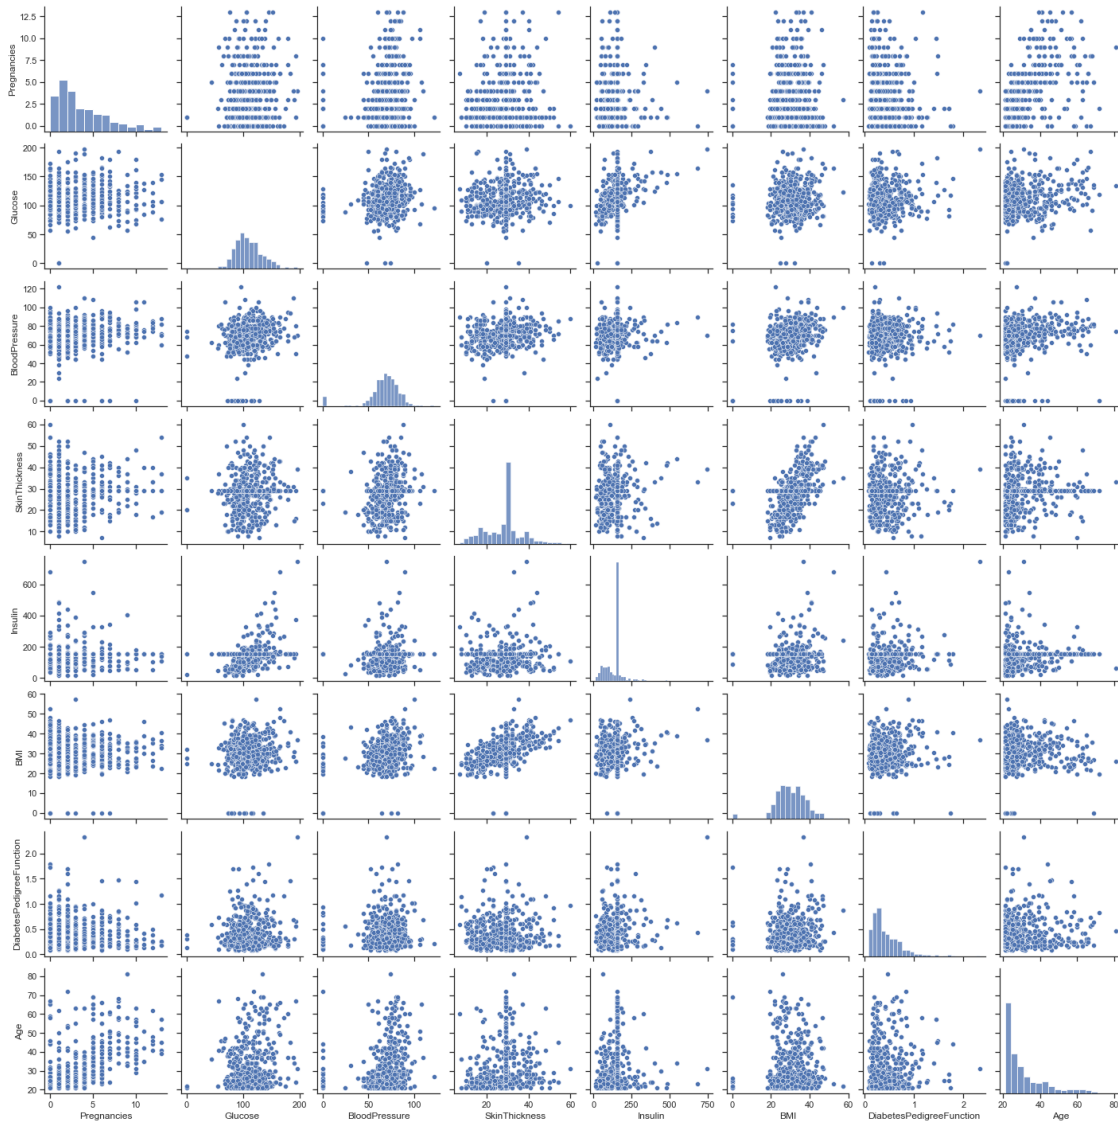
```
[253]: #Pair plots for all dataset
      sns.set(style="ticks", color_codes=True)
      g = sns.pairplot(dataset_imputed, hue="Outcome")
```



```
[254]: #Pair plots for all Positive cases
sns.set(style="ticks", color_codes=True)
g = sns.
    ↪ pairplot(Positive[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
    ↪ 'Age']])
```



```
[255]: #Pair plots for all Negative cases
sns.set(style="ticks", color_codes=True)
g = sns.
    ↪ pairplot(Negative[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
    ↪ 'Age']])
```



0.1.3 Correlation Analysis and Heat map

```
[256]: ### correlation matrix
dataset_imputed.corr()
```

```
[256]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	0.082989	
Glucose	0.129459	1.000000	0.152590	0.182455	
BloodPressure	0.141282	0.152590	1.000000	0.123444	
SkinThickness	0.082989	0.182455	0.123444	1.000000	
Insulin	0.056027	0.407699	0.045319	0.158139	
BMI	0.017683	0.221071	0.281805	0.480496	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.100966	


Age	0.544341	0.263514	0.239528	0.127872
Outcome	0.221898	0.466581	0.065068	0.215299

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	0.056027	0.017683	-0.033523	
Glucose	0.407699	0.221071	0.137337	
BloodPressure	0.045319	0.281805	0.041265	
SkinThickness	0.158139	0.480496	0.100966	
Insulin	1.000000	0.149468	0.098634	
BMI	0.149468	1.000000	0.140647	
DiabetesPedigreeFunction	0.098634	0.140647	1.000000	
Age	0.136734	0.036242	0.033561	
Outcome	0.214411	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	0.127872	0.215299
Insulin	0.136734	0.214411
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[257]: plt.subplots(figsize=(8,8))
sns.heatmap(dataset_imputed.corr(),annot=True)
```

```
[257]: <AxesSubplot:>
```



0.1.4 Project Task: Week 3 and Week 4 – Data Modelling and Model Performance Evaluation

Model 1 : Logistic Regression

```
[258]: dataset_imputed.head(5)
```

```
[258]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6      148           72      35.00000  155.548223  33.6
1           1       85           66      29.00000  155.548223  26.6
2           8      183           64      29.15342  155.548223  23.3
3           1       89           66      23.00000   94.000000  28.1
4           0      137           40      35.00000  168.000000  43.1
```

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[259]: features = dataset_imputed.iloc[:,[0,1,2,3,4,5,6,7]].values
label = dataset_imputed.iloc[:,8].values
```

```
[260]: #Train test split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(features,
                                                label,
                                                test_size=0.2,
                                                random_state =10)
```

```
[261]: #Create model
from sklearn.linear_model import LogisticRegression
logRegModel = LogisticRegression()
logRegModel.fit(X_train,y_train)
```

```
[261]: LogisticRegression()
```

```
[308]: y_pred = logRegModel.predict(X_test)
from sklearn.metrics import accuracy_score
print('Accuracy of logistic regression classifier on test set',
      accuracy_score(y_test, y_pred))
```

Accuracy of logistic regression classifier on test set 0.7597402597402597

```
[309]: from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[86  9]
 [28 31]]
```

Model 2 : Decision Tree Classifier

```
[310]: #Hyper Parameter tuning of max_dept
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
for i in range(3,20):
    print("For max_depth = ",i)
    DTModel = DecisionTreeClassifier(max_depth=i)
    DTModel.fit(X_train,y_train)
```



```
y_pred = DTModel.predict(X_test)
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

```
For max_depth = 3
Accuracy: 0.6948051948051948
For max_depth = 4
Accuracy: 0.7532467532467533
For max_depth = 5
Accuracy: 0.7597402597402597
For max_depth = 6
Accuracy: 0.7467532467532467
For max_depth = 7
Accuracy: 0.7597402597402597
For max_depth = 8
Accuracy: 0.7532467532467533
For max_depth = 9
Accuracy: 0.7727272727272727
For max_depth = 10
Accuracy: 0.7727272727272727
For max_depth = 11
Accuracy: 0.7337662337662337
For max_depth = 12
Accuracy: 0.7012987012987013
For max_depth = 13
Accuracy: 0.7012987012987013
For max_depth = 14
Accuracy: 0.7402597402597403
For max_depth = 15
Accuracy: 0.6948051948051948
For max_depth = 16
Accuracy: 0.7142857142857143
For max_depth = 17
Accuracy: 0.6818181818181818
For max_depth = 18
Accuracy: 0.7142857142857143
For max_depth = 19
Accuracy: 0.7272727272727273
```

Highest Accuracy of Decision Tree Model can be obtained on Max_Depth = 10

```
[311]: DTModel = DecisionTreeClassifier(max_depth=10)
DTModel.fit(X_train,y_train)
y_pred = DTModel.predict(X_test)
```

```
[312]: DTModel.score(X_train,y_train)
```

```
[312]: 0.9267100977198697
```

```
[313]: DTModel.score(X_test,y_test)
```

```
[313]: 0.7532467532467533
```

```
[315]: print('Accuracy of Decision Tree regression classifier on test set',  
        ↪accuracy_score(y_test, y_pred))
```

Accuracy of Decision Tree regression classifier on test set 0.7532467532467533

```
[272]: from sklearn.metrics import confusion_matrix  
confusion_matrix = confusion_matrix(y_test, y_pred)  
print(confusion_matrix)
```

```
[[77 18]  
 [20 39]]
```

Model 3 : Random Forest Classifier

```
[277]: from sklearn.ensemble import RandomForestClassifier  
rf = RandomForestClassifier()  
rf.fit(X_train, y_train)  
y_pred = rf.predict(X_test)
```

```
[281]: rfModel = RandomForestClassifier(n_estimators=60)  
rfModel.fit(X_train, y_train)  
y_pred = rfModel.predict(X_test)
```

```
[317]: print('Accuracy of Random Forest regression classifier on test set',  
        ↪accuracy_score(y_test, y_pred))
```

Accuracy of Random Forest regression classifier on test set 0.7532467532467533

```
[286]: from sklearn.metrics import confusion_matrix  
confusion_matrix = confusion_matrix(y_test, y_pred)  
print(confusion_matrix)
```

```
[[85 10]  
 [27 32]]
```

Model 4 : Support Vector Machine

```
[291]: #Support Vector Classifier  
  
from sklearn.svm import SVC  
SVMmodel = SVC(kernel='rbf',  
                gamma='auto')  
SVMmodel.fit(X_train,y_train)
```

```
[291]: SVC(gamma='auto')
```

```
[318]: y_pred=SVMmodel.predict(X_test)
```

```
[319]: print('Accuracy of Support Vector Machine on test set', accuracy_score(y_test, y_pred))
```

Accuracy of Support Vector Machine on test set 0.6168831168831169

Model 5 : KNN Classifier

```
[294]: #Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
knnClassifier = KNeighborsClassifier(n_neighbors=7,
                                    metric='minkowski',
                                    p = 2)
knnClassifier.fit(X_train,y_train)
```

```
[294]: KNeighborsClassifier(n_neighbors=7)
```

```
[320]: y_pred=knnClassifier.predict(X_test)
```

```
[321]: print('Accuracy of KNN Classifier on test set', accuracy_score(y_test, y_pred))
```

Accuracy of KNN Classifier on test set 0.7272727272727273

We observed that Random Forest is best performing model for this dataset

Accuracy of 75%

```
[ ]:
```